



```
from google.colab import files
files.upload()
```

 no files selected Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving kaggle.json to kaggle.json
{'kaggle.json':

```
!mkdir ~/.kaggle/
```

```
!cp kaggle.json ~/.kaggle/  
!chmod 600 ~/.kaggle/kaggle.json
```

```
!kaggle competitions download -c dogs-vs-cats
```

 Downloading dogs-vs-cats.zip to /content
99% 803M/812M [00:09<00:00, 61.4MB/s]
100% 812M/812M [00:09<00:00, 89.3MB/s]

```
!unzip -qq dogs-vs-cats.zip
```

```
!unzip -qq train.zip
```

1.Consider the Cats & Dogs example. Start initially with a training sample of 1000, a validation sample of 500, and a test sample of 500 (like in the text). Use any technique to reduce overfitting and improve performance in developing a network that you train from scratch. What performance did you achieve?

Copying pictures to the test, validation, and training sets.

```
import os, shutil, pathlib

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small_1")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index, end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=1000)
make_subset("validation", start_index=1000, end_index=1500)
make_subset("test", start_index=1500, end_index=2000)
```

Convolutional neural networks are used.

```
from tensorflow import keras
from tensorflow.keras import layers

inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
```

```
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 180, 180, 3)]	0
rescaling (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
conv2d_2 (Conv2D)	(None, 41, 41, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 20, 20, 128)	0
conv2d_3 (Conv2D)	(None, 18, 18, 256)	295168
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 256)	0
conv2d_4 (Conv2D)	(None, 7, 7, 256)	590080
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 1)	12545

```

Total params: 991041 (3.78 MB)
Trainable params: 991041 (3.78 MB)
Non-trainable params: 0 (0.00 Byte)

```

Given that the model may overfit, regularization techniques are employed at the DATA PREPROCESSING step.

Here, every image is transformed into a tensor.

```
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
```

```
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)
```

```
Found 2000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
```

Callback can be either used for saveage the model's weights at the end of every epoch or for early stopping if the model is not improving. Besides that callbacks, such as logging metrics, visualizing the model performance, or scheduling a learning rate changes may be used in the same way.

```
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch1.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)
```

Epoch 1/30

```
63/63 [=====] - 6s 96ms/step - loss: 0.6868 - accuracy: 0.7500  
Epoch 2/30  
63/63 [=====] - 4s 57ms/step - loss: 0.6487 - accuracy: 0.7500  
Epoch 3/30  
63/63 [=====] - 4s 56ms/step - loss: 0.5969 - accuracy: 0.7500  
Epoch 4/30  
63/63 [=====] - 5s 72ms/step - loss: 0.5681 - accuracy: 0.7500  
Epoch 5/30  
63/63 [=====] - 6s 84ms/step - loss: 0.5399 - accuracy: 0.7500  
Epoch 6/30  
63/63 [=====] - 4s 57ms/step - loss: 0.5165 - accuracy: 0.7500  
Epoch 7/30  
63/63 [=====] - 7s 103ms/step - loss: 0.4856 - accuracy: 0.7500  
Epoch 8/30  
63/63 [=====] - 4s 56ms/step - loss: 0.4560 - accuracy: 0.7500  
Epoch 9/30  
63/63 [=====] - 4s 56ms/step - loss: 0.4115 - accuracy: 0.7500  
Epoch 10/30  
63/63 [=====] - 6s 97ms/step - loss: 0.3636 - accuracy: 0.7500  
Epoch 11/30  
63/63 [=====] - 4s 57ms/step - loss: 0.3278 - accuracy: 0.7500  
Epoch 12/30  
63/63 [=====] - 4s 60ms/step - loss: 0.2721 - accuracy: 0.7500  
Epoch 13/30  
63/63 [=====] - 6s 83ms/step - loss: 0.2251 - accuracy: 0.7500  
Epoch 14/30  
63/63 [=====] - 4s 58ms/step - loss: 0.1930 - accuracy: 0.7500  
Epoch 15/30  
63/63 [=====] - 4s 57ms/step - loss: 0.1472 - accuracy: 0.7500  
Epoch 16/30  
63/63 [=====] - 5s 80ms/step - loss: 0.1295 - accuracy: 0.7500  
Epoch 17/30  
63/63 [=====] - 4s 56ms/step - loss: 0.0976 - accuracy: 0.7500  
Epoch 18/30  
63/63 [=====] - 6s 90ms/step - loss: 0.0945 - accuracy: 0.7500  
Epoch 19/30  
63/63 [=====] - 4s 57ms/step - loss: 0.0674 - accuracy: 0.7500  
Epoch 20/30  
63/63 [=====] - 4s 56ms/step - loss: 0.0546 - accuracy: 0.7500  
Epoch 21/30  
63/63 [=====] - 6s 87ms/step - loss: 0.0719 - accuracy: 0.7500  
Epoch 22/30  
63/63 [=====] - 4s 57ms/step - loss: 0.0767 - accuracy: 0.7500  
Epoch 23/30  
63/63 [=====] - 6s 95ms/step - loss: 0.0494 - accuracy: 0.7500  
Epoch 24/30  
63/63 [=====] - 5s 66ms/step - loss: 0.0561 - accuracy: 0.7500  
Epoch 25/30  
63/63 [=====] - 4s 61ms/step - loss: 0.0406 - accuracy: 0.7500  
Epoch 26/30
```

```

63/63 [=====] - 4s 64ms/step - loss: 0.0462 - accuracy: 0.713
Epoch 27/30
63/63 [=====] - 7s 103ms/step - loss: 0.0544 - accuracy: 0.704
Epoch 28/30
63/63 [=====] - 4s 60ms/step - loss: 0.0533 - accuracy: 0.713
Epoch 29/30
63/63 [=====] - 6s 95ms/step - loss: 0.0435 - accuracy: 0.713
Epoch 30/30

```


Accuracy appears to be rising with the number of epochs.


Accuracy=71.3 Val_acc=70.4 test_acc=71.3

```

import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("training and Validation Accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("training and Validation Loss")
plt.legend()
plt.show()

```





```
test_model = keras.models.load_model("convnet_from_scratch1.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
```

```
32/32 [=====] - 1s 28ms/step - loss: 0.5793 - accuracy: 0.713
Test accuracy: 0.713
```

Test accuracy with no data augmentation=71.3

✓ Data Augmentation

Data augmentation acts like a master artist adding new custom made artwork to the training dataset by artfully modifying the original data. Apart from assisting in curbing overfitting, this aesthetic touch also leads to the development of a broader capability of recognizing more generalized cases by the model.

```
data_augmentation = keras.Sequential(  
    [  
        layers.RandomFlip("horizontal"),  
        layers.RandomRotation(0.1),  
        layers.RandomZoom(0.2),  
    ]  
)
```

```
inputs = keras.Input(shape=(180, 180, 3))  
x = data_augmentation(inputs)  
x = layers.Rescaling(1./255)(x)  
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)  
x = layers.MaxPooling2D(pool_size=2)(x)  
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)  
x = layers.Flatten()(x)  
x = layers.Dropout(0.5)(x)  
outputs = layers.Dense(1, activation="sigmoid")(x)  
model = keras.Model(inputs=inputs, outputs=outputs)  
  
model.compile(loss="binary_crossentropy",  
              optimizer="rmsprop",  
              metrics=["accuracy"])
```

```
callbacks = [  
    keras.callbacks.ModelCheckpoint(  
        filepath="convnet_from_scratch_with_augmentation1.keras",
```



```

        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=50,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

```

Epoch 1/50
63/63 [=====] - 9s 95ms/step - loss: 0.6988 - accuracy: 0.1250
Epoch 2/50
63/63 [=====] - 5s 75ms/step - loss: 0.6943 - accuracy: 0.1250
Epoch 3/50
63/63 [=====] - 6s 85ms/step - loss: 0.6910 - accuracy: 0.1250
Epoch 4/50
63/63 [=====] - 6s 83ms/step - loss: 0.6906 - accuracy: 0.1250
Epoch 5/50
63/63 [=====] - 4s 62ms/step - loss: 0.6798 - accuracy: 0.1250
Epoch 6/50
63/63 [=====] - 6s 93ms/step - loss: 0.6610 - accuracy: 0.1250
Epoch 7/50
63/63 [=====] - 4s 59ms/step - loss: 0.6469 - accuracy: 0.1250
Epoch 8/50
63/63 [=====] - 4s 60ms/step - loss: 0.6402 - accuracy: 0.1250
Epoch 9/50
63/63 [=====] - 8s 114ms/step - loss: 0.6243 - accuracy: 0.1250
Epoch 10/50
63/63 [=====] - 5s 67ms/step - loss: 0.6165 - accuracy: 0.1250
Epoch 11/50
63/63 [=====] - 6s 86ms/step - loss: 0.5941 - accuracy: 0.1250
Epoch 12/50
63/63 [=====] - 4s 60ms/step - loss: 0.5913 - accuracy: 0.1250
Epoch 13/50
63/63 [=====] - 6s 86ms/step - loss: 0.5724 - accuracy: 0.1250
Epoch 14/50
63/63 [=====] - 6s 85ms/step - loss: 0.5659 - accuracy: 0.1250
Epoch 15/50
63/63 [=====] - 4s 60ms/step - loss: 0.5640 - accuracy: 0.1250
Epoch 16/50
63/63 [=====] - 4s 62ms/step - loss: 0.5350 - accuracy: 0.1250
Epoch 17/50
63/63 [=====] - 6s 90ms/step - loss: 0.5368 - accuracy: 0.1250
Epoch 18/50
63/63 [=====] - 4s 58ms/step - loss: 0.5253 - accuracy: 0.1250
Epoch 19/50
63/63 [=====] - 4s 62ms/step - loss: 0.5116 - accuracy: 0.1250
Epoch 20/50

```

```
63/63 [=====] - 6s 90ms/step - loss: 0.5123 - accuracy: 0.7500  
Epoch 21/50  
63/63 [=====] - 4s 59ms/step - loss: 0.4917 - accuracy: 0.7500  
Epoch 22/50  
63/63 [=====] - 7s 110ms/step - loss: 0.4862 - accuracy: 0.7500  
Epoch 23/50  
63/63 [=====] - 4s 58ms/step - loss: 0.4734 - accuracy: 0.7500  
Epoch 24/50  
63/63 [=====] - 5s 71ms/step - loss: 0.4826 - accuracy: 0.7500  
Epoch 25/50  
63/63 [=====] - 6s 91ms/step - loss: 0.4590 - accuracy: 0.7500  
Epoch 26/50  
63/63 [=====] - 4s 58ms/step - loss: 0.4477 - accuracy: 0.7500  
Epoch 27/50  
63/63 [=====] - 7s 101ms/step - loss: 0.4377 - accuracy: 0.7500  
Epoch 28/50  
63/63 [=====] - 4s 62ms/step - loss: 0.4309 - accuracy: 0.7500  
Epoch 29/50  
63/63 [=====] - 4s 59ms/step - loss: 0.4420 - accuracy: 0.7500  
Epoch 30/50
```

Displaying some randomly augmented training images:

```
plt.figure(figsize=(7.5,7.5 ))
for images, _ in train_dataset.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(augmented_images[0].numpy().astype("uint8"))
        plt.axis("off")
```



2. Increase your training sample size. You may pick any amount. Keep the validation and test samples the same as above

Tried to raise the 1000–1500 training sample size.

```
import os, shutil, pathlib

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small_2")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir, exist_ok=True)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index, end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=1500)
make_subset("validation", start_index=1500, end_index=2000)
make_subset("test", start_index=2000, end_index=2500)
```

```
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)
```

```
Found 3000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
```

```

inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch2.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=70,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

```

Epoch 1/70
94/94 [=====] - 9s 78ms/step - loss: 0.6923 - accuracy: 0.2500
Epoch 2/70
94/94 [=====] - 7s 69ms/step - loss: 0.6699 - accuracy: 0.2500
Epoch 3/70
94/94 [=====] - 5s 52ms/step - loss: 0.6373 - accuracy: 0.2500
Epoch 4/70
94/94 [=====] - 8s 78ms/step - loss: 0.5955 - accuracy: 0.2500
Epoch 5/70
94/94 [=====] - 5s 52ms/step - loss: 0.5593 - accuracy: 0.2500
Epoch 6/70
94/94 [=====] - 7s 71ms/step - loss: 0.5061 - accuracy: 0.2500
Epoch 7/70

```

```
94/94 [=====] - 6s 64ms/step - loss: 0.4537 - accuracy: 0.7500
Epoch 8/70
94/94 [=====] - 5s 51ms/step - loss: 0.4003 - accuracy: 0.7500
Epoch 9/70
94/94 [=====] - 7s 75ms/step - loss: 0.3641 - accuracy: 0.7500
Epoch 10/70
94/94 [=====] - 6s 61ms/step - loss: 0.2999 - accuracy: 0.7500
Epoch 11/70
94/94 [=====] - 6s 57ms/step - loss: 0.2355 - accuracy: 0.7500
Epoch 12/70
94/94 [=====] - 9s 93ms/step - loss: 0.1936 - accuracy: 0.7500
Epoch 13/70
94/94 [=====] - 5s 54ms/step - loss: 0.1570 - accuracy: 0.7500
Epoch 14/70
94/94 [=====] - 7s 70ms/step - loss: 0.1067 - accuracy: 0.7500
Epoch 15/70
94/94 [=====] - 5s 51ms/step - loss: 0.0928 - accuracy: 0.7500
Epoch 16/70
94/94 [=====] - 7s 72ms/step - loss: 0.0842 - accuracy: 0.7500
Epoch 17/70
94/94 [=====] - 7s 64ms/step - loss: 0.0665 - accuracy: 0.7500
Epoch 18/70
94/94 [=====] - 7s 73ms/step - loss: 0.0608 - accuracy: 0.7500
Epoch 19/70
94/94 [=====] - 6s 63ms/step - loss: 0.0555 - accuracy: 0.7500
Epoch 20/70
94/94 [=====] - 5s 51ms/step - loss: 0.0509 - accuracy: 0.7500
Epoch 21/70
94/94 [=====] - 8s 83ms/step - loss: 0.0356 - accuracy: 0.7500
Epoch 22/70
94/94 [=====] - 5s 52ms/step - loss: 0.0377 - accuracy: 0.7500
Epoch 23/70
94/94 [=====] - 9s 92ms/step - loss: 0.0521 - accuracy: 0.7500
Epoch 24/70
94/94 [=====] - 5s 51ms/step - loss: 0.0394 - accuracy: 0.7500
Epoch 25/70
94/94 [=====] - 6s 57ms/step - loss: 0.0336 - accuracy: 0.7500
Epoch 26/70
94/94 [=====] - 8s 77ms/step - loss: 0.0472 - accuracy: 0.7500
Epoch 27/70
94/94 [=====] - 5s 51ms/step - loss: 0.0496 - accuracy: 0.7500
Epoch 28/70
94/94 [=====] - 7s 69ms/step - loss: 0.0310 - accuracy: 0.7500
Epoch 29/70
94/94 [=====] - 7s 70ms/step - loss: 0.0512 - accuracy: 0.7500
Epoch 30/70
```

```
test_model = keras.models.load_model(
    "convnet_from_scratch2.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
```

```
32/32 [=====] - 2s 33ms/step - loss: 0.5028 - accuracy: 0.758
Test accuracy: 0.758
```

Accuracy=75.8 val_acc=74.6 test_acc=75.8

✓ Applying data augmentation

```
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)
```

```

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="adam",
              metrics=["accuracy"])

```

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation2.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=80,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

```

Epoch 1/80
94/94 [=====] - 6s 57ms/step - loss: 0.6896 - accuracy: 0.7500
Epoch 2/80
94/94 [=====] - 8s 81ms/step - loss: 0.6922 - accuracy: 0.7500
Epoch 3/80
94/94 [=====] - 6s 56ms/step - loss: 0.6827 - accuracy: 0.7500
Epoch 4/80
94/94 [=====] - 9s 93ms/step - loss: 0.6571 - accuracy: 0.7500
Epoch 5/80
94/94 [=====] - 5s 55ms/step - loss: 0.6494 - accuracy: 0.7500
Epoch 6/80

```



```

94/94 [=====] - 7s 68ms/step - loss: 0.6400 - accuracy: 0.7700
Epoch 7/80
94/94 [=====] - 5s 53ms/step - loss: 0.6266 - accuracy: 0.7700
Epoch 8/80
94/94 [=====] - 7s 72ms/step - loss: 0.5966 - accuracy: 0.7700
Epoch 9/80
94/94 [=====] - 5s 53ms/step - loss: 0.5968 - accuracy: 0.7700
Epoch 10/80
94/94 [=====] - 7s 69ms/step - loss: 0.5732 - accuracy: 0.7700
Epoch 11/80
94/94 [=====] - 7s 72ms/step - loss: 0.5598 - accuracy: 0.7700
Epoch 12/80
94/94 [=====] - 6s 63ms/step - loss: 0.5337 - accuracy: 0.7700
Epoch 13/80
94/94 [=====] - 5s 53ms/step - loss: 0.5350 - accuracy: 0.7700
Epoch 14/80
94/94 [=====] - 6s 56ms/step - loss: 0.4995 - accuracy: 0.7700
Epoch 15/80
94/94 [=====] - 5s 52ms/step - loss: 0.5053 - accuracy: 0.7700
Epoch 16/80
94/94 [=====] - 8s 85ms/step - loss: 0.4901 - accuracy: 0.7700
Epoch 17/80
94/94 [=====] - 5s 52ms/step - loss: 0.4780 - accuracy: 0.7700
Epoch 18/80
94/94 [=====] - 8s 78ms/step - loss: 0.4719 - accuracy: 0.7700
Epoch 19/80
94/94 [=====] - 5s 56ms/step - loss: 0.4754 - accuracy: 0.7700
Epoch 20/80
94/94 [=====] - 7s 72ms/step - loss: 0.4445 - accuracy: 0.7700
Epoch 21/80
94/94 [=====] - 6s 63ms/step - loss: 0.4477 - accuracy: 0.7700
Epoch 22/80
94/94 [=====] - 7s 68ms/step - loss: 0.4319 - accuracy: 0.7700
Epoch 23/80
94/94 [=====] - 5s 53ms/step - loss: 0.4268 - accuracy: 0.7700
Epoch 24/80
94/94 [=====] - 6s 64ms/step - loss: 0.4227 - accuracy: 0.7700
Epoch 25/80
94/94 [=====] - 5s 52ms/step - loss: 0.4176 - accuracy: 0.7700
Epoch 26/80
94/94 [=====] - 7s 72ms/step - loss: 0.3990 - accuracy: 0.7700
Epoch 27/80
94/94 [=====] - 5s 53ms/step - loss: 0.3953 - accuracy: 0.7700
Epoch 28/80
94/94 [=====] - 7s 71ms/step - loss: 0.3984 - accuracy: 0.7700
Epoch 29/80
94/94 [=====] - 7s 68ms/step - loss: 0.3863 - accuracy: 0.7700
Epoch 30/80
94/94 [=====] - 7s 72ms/step - loss: 0.3760 - accuracy: 0.7700

```

Accuracy=93.8 val_acc=85.2 test_acc=84.8

```
test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation2.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
```

```
32/32 [=====] - 2s 36ms/step - loss: 0.3820 - accuracy: 0.8480
Test accuracy: 0.848
```

3. The objective is to find the ideal training sample size to get best prediction

The sizes of the training, validation, and test sets were established at 1500, 1000, and 500, respectively.

```
import os, shutil, pathlib

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small_3")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir, exist_ok=True)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index, end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=1500)
make_subset("validation", start_index=1500, end_index=2500)
make_subset("test", start_index=2500, end_index=3000)
```

```
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)
```

```
Found 3000 files belonging to 2 classes.
Found 2000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
```

```
inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
```

```
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch3.keras",
        save_best_only=True,
        monitor="val_loss")]
```

```
]
history = model.fit(
    train_dataset,
    epochs=90,
    validation_data=validation_dataset,
    callbacks=callbacks)
```

```
Epoch 1/90
94/94 [=====] - 7s 72ms/step - loss: 0.6906 - accuracy: 0.1250
Epoch 2/90
94/94 [=====] - 6s 61ms/step - loss: 0.6559 - accuracy: 0.1562
Epoch 3/90
94/94 [=====] - 8s 78ms/step - loss: 0.6158 - accuracy: 0.1875
Epoch 4/90
94/94 [=====] - 7s 67ms/step - loss: 0.5944 - accuracy: 0.2187
Epoch 5/90
94/94 [=====] - 8s 84ms/step - loss: 0.5553 - accuracy: 0.2500
Epoch 6/90
94/94 [=====] - 6s 61ms/step - loss: 0.5029 - accuracy: 0.2812
Epoch 7/90
94/94 [=====] - 8s 80ms/step - loss: 0.4587 - accuracy: 0.3125
Epoch 8/90
94/94 [=====] - 6s 60ms/step - loss: 0.4258 - accuracy: 0.3437
Epoch 9/90
94/94 [=====] - 9s 98ms/step - loss: 0.3608 - accuracy: 0.3750
Epoch 10/90
94/94 [=====] - 7s 67ms/step - loss: 0.3092 - accuracy: 0.4062
Epoch 11/90
94/94 [=====] - 8s 86ms/step - loss: 0.2478 - accuracy: 0.4375
Epoch 12/90
94/94 [=====] - 6s 61ms/step - loss: 0.2054 - accuracy: 0.4687
Epoch 13/90
94/94 [=====] - 8s 83ms/step - loss: 0.1628 - accuracy: 0.5000
Epoch 14/90
94/94 [=====] - 7s 69ms/step - loss: 0.1234 - accuracy: 0.5312
Epoch 15/90
94/94 [=====] - 6s 65ms/step - loss: 0.0944 - accuracy: 0.5625
Epoch 16/90
94/94 [=====] - 9s 89ms/step - loss: 0.0906 - accuracy: 0.5937
Epoch 17/90
94/94 [=====] - 7s 67ms/step - loss: 0.0660 - accuracy: 0.6250
Epoch 18/90
94/94 [=====] - 9s 94ms/step - loss: 0.0746 - accuracy: 0.6562
Epoch 19/90
94/94 [=====] - 6s 60ms/step - loss: 0.0463 - accuracy: 0.6875
Epoch 20/90
94/94 [=====] - 9s 89ms/step - loss: 0.0588 - accuracy: 0.7187
Epoch 21/90
```

```

94/94 [=====] - 7s 67ms/step - loss: 0.0366 - accuracy: 0.993
Epoch 22/90
94/94 [=====] - 7s 67ms/step - loss: 0.0454 - accuracy: 0.993
Epoch 23/90
94/94 [=====] - 9s 93ms/step - loss: 0.0344 - accuracy: 0.993
Epoch 24/90
94/94 [=====] - 6s 61ms/step - loss: 0.0464 - accuracy: 0.993
Epoch 25/90
94/94 [=====] - 7s 74ms/step - loss: 0.0591 - accuracy: 0.993
Epoch 26/90
94/94 [=====] - 8s 79ms/step - loss: 0.0348 - accuracy: 0.993
Epoch 27/90
94/94 [=====] - 9s 96ms/step - loss: 0.0388 - accuracy: 0.993
Epoch 28/90
94/94 [=====] - 6s 60ms/step - loss: 0.0358 - accuracy: 0.993
Epoch 29/90
94/94 [=====] - 8s 78ms/step - loss: 0.0654 - accuracy: 0.993
Epoch 30/90

```

```

test_model = keras.models.load_model(
    "convnet_from_scratch3.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

```

```

32/32 [=====] - 1s 28ms/step - loss: 0.5366 - accuracy: 0.752
Test accuracy: 0.752

```

Accuracy=99.3 val_Acc=75.0 test_Acc=75.2

✓ Utilizing Augmented Data

```

data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)

```

```

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="adam",
              metrics=["accuracy"])

```

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation3.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=100,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

```

Epoch 1/100
94/94 [=====] - 13s 107ms/step - loss: 0.6933 - accur
Epoch 2/100
94/94 [=====] - 9s 94ms/step - loss: 0.6943 - accur
Epoch 3/100
94/94 [=====] - 6s 62ms/step - loss: 0.6941 - accur
Epoch 4/100
94/94 [=====] - 7s 71ms/step - loss: 0.6935 - accur
Epoch 5/100
94/94 [=====] - 9s 92ms/step - loss: 0.6934 - accur
Epoch 6/100
94/94 [=====] - 6s 64ms/step - loss: 0.6933 - accur

```

```

Epoch 7/100
94/94 [=====] - 10s 104ms/step - loss: 0.6934 - accu
Epoch 8/100
94/94 [=====] - 6s 61ms/step - loss: 0.6934 - accur
Epoch 9/100
94/94 [=====] - 9s 89ms/step - loss: 0.6958 - accur
Epoch 10/100
94/94 [=====] - 6s 62ms/step - loss: 0.6934 - accur
Epoch 11/100
94/94 [=====] - 7s 69ms/step - loss: 0.6921 - accur
Epoch 12/100
94/94 [=====] - 10s 109ms/step - loss: 0.6945 - accu
Epoch 13/100
94/94 [=====] - 6s 62ms/step - loss: 0.6933 - accur
Epoch 14/100
94/94 [=====] - 8s 84ms/step - loss: 0.6933 - accur
Epoch 15/100
94/94 [=====] - 9s 96ms/step - loss: 0.6932 - accur
Epoch 16/100
94/94 [=====] - 7s 68ms/step - loss: 0.6927 - accur
Epoch 17/100
94/94 [=====] - 9s 96ms/step - loss: 0.6932 - accur
Epoch 18/100
94/94 [=====] - 7s 66ms/step - loss: 0.6934 - accur
Epoch 19/100
94/94 [=====] - 7s 71ms/step - loss: 0.6931 - accur
Epoch 20/100
94/94 [=====] - 9s 86ms/step - loss: 0.6938 - accur
Epoch 21/100
94/94 [=====] - 6s 61ms/step - loss: 0.6929 - accur
Epoch 22/100
94/94 [=====] - 12s 122ms/step - loss: 0.6932 - accu
Epoch 23/100
94/94 [=====] - 6s 61ms/step - loss: 0.6930 - accur
Epoch 24/100
94/94 [=====] - 7s 65ms/step - loss: 0.6930 - accur
Epoch 25/100
94/94 [=====] - 9s 96ms/step - loss: 0.6929 - accur
Epoch 26/100
94/94 [=====] - 6s 62ms/step - loss: 0.6942 - accur
Epoch 27/100
94/94 [=====] - 8s 81ms/step - loss: 0.6933 - accur
Epoch 28/100
94/94 [=====] - 8s 75ms/step - loss: 0.6926 - accur
Epoch 29/100
94/94 [=====] - 7s 68ms/step - loss: 0.6936 - accur
Epoch 30/100
- - - - -

```

```
test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation3.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
```

```
32/32 [=====] - 1s 28ms/step - loss: 0.6919 - accuracy: 0.524
Test accuracy: 0.524
```

Accuracy=48.9 val_acc=50.0 test_acc=52.4

✓ 4. Now using a pretrained network.

This pre-trained network is modeled by VGG16.

Creating the VGG16 convolutional basis is the first step in feature extraction.

```
conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False,
    input_shape=(180, 180, 3))
conv_base.summary()
```


Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/58889256/58889256> [=====] - 0s 0us/step
 Model: "vgg16"

Layer (type)	Output Shape	Param #
input_7 (InputLayer)	[(None, 180, 180, 3)]	0
block1_conv1 (Conv2D)	(None, 180, 180, 64)	1792
block1_conv2 (Conv2D)	(None, 180, 180, 64)	36928
block1_pool (MaxPooling2D)	(None, 90, 90, 64)	0
block2_conv1 (Conv2D)	(None, 90, 90, 128)	73856
block2_conv2 (Conv2D)	(None, 90, 90, 128)	147584
block2_pool (MaxPooling2D)	(None, 45, 45, 128)	0
block3_conv1 (Conv2D)	(None, 45, 45, 256)	295168
block3_conv2 (Conv2D)	(None, 45, 45, 256)	590080
block3_conv3 (Conv2D)	(None, 45, 45, 256)	590080
block3_pool (MaxPooling2D)	(None, 22, 22, 256)	0
block4_conv1 (Conv2D)	(None, 22, 22, 512)	1180160
block4_conv2 (Conv2D)	(None, 22, 22, 512)	2359808
block4_conv3 (Conv2D)	(None, 22, 22, 512)	2359808
block4_pool (MaxPooling2D)	(None, 11, 11, 512)	0
block5_conv1 (Conv2D)	(None, 11, 11, 512)	2359808
block5_conv2 (Conv2D)	(None, 11, 11, 512)	2359808
block5_conv3 (Conv2D)	(None, 11, 11, 512)	2359808
block5_pool (MaxPooling2D)	(None, 5, 5, 512)	0
=====		
Total params: 14714688 (56.13 MB)		
Trainable params: 14714688 (56.13 MB)		
Non-trainable params: 0 (0.00 Byte)		

Extraction of features and associated labels is known as feature extraction.

```
import numpy as np

def get_features_and_labels(dataset):
    all_features = []
    all_labels = []
    for images, labels in dataset:
        preprocessed_images = keras.applications.vgg16.preprocess_input(images)
        features = conv_base.predict(preprocessed_images)
        all_features.append(features)
        all_labels.append(labels)
    return np.concatenate(all_features), np.concatenate(all_labels)

train_features, train_labels = get_features_and_labels(train_dataset)
val_features, val_labels = get_features_and_labels(validation_dataset)
test_features, test_labels = get_features_and_labels(test_dataset)

train_features.shape
```

```
1/1 [=====] - 5s 5s/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 45ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 49ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 43ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 33ms/step
```

```

1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step

```

The densely connected classifier is defined and trained throughout the feature extraction process.

```

inputs = keras.Input(shape=(5, 5, 512))
x = layers.Flatten()(inputs)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extractionPT1.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_features, train_labels,
    epochs=15,
    validation_data=(val_features, val_labels),
    callbacks=callbacks)

```


```


Epoch 1/15
94/94 [=====] - 2s 10ms/step - loss: 19.6264 - accuracy: 0.0000
Epoch 2/15
94/94 [=====] - 0s 5ms/step - loss: 3.1697 - accuracy: 0.0000
Epoch 3/15
94/94 [=====] - 0s 5ms/step - loss: 1.6974 - accuracy: 0.0000
Epoch 4/15
94/94 [=====] - 0s 5ms/step - loss: 0.9081 - accuracy: 0.0000
Epoch 5/15
94/94 [=====] - 0s 5ms/step - loss: 0.4778 - accuracy: 0.0000
Epoch 6/15
94/94 [=====] - 1s 8ms/step - loss: 1.3842 - accuracy: 0.0000
Epoch 7/15
94/94 [=====] - 1s 8ms/step - loss: 0.5023 - accuracy: 0.0000
Epoch 8/15
94/94 [=====] - 1s 8ms/step - loss: 0.4250 - accuracy: 0.0000
Epoch 9/15
94/94 [=====] - 1s 8ms/step - loss: 0.3422 - accuracy: 0.0000
Epoch 10/15
94/94 [=====] - 1s 8ms/step - loss: 0.4261 - accuracy: 0.0000
Epoch 11/15
94/94 [=====] - 1s 8ms/step - loss: 0.3498 - accuracy: 0.0000
Epoch 12/15
94/94 [=====] - 1s 8ms/step - loss: 0.2699 - accuracy: 0.0000
Epoch 13/15
94/94 [=====] - 1s 9ms/step - loss: 0.0019 - accuracy: 0.0000
Epoch 14/15
94/94 [=====] - 1s 7ms/step - loss: 0.0149 - accuracy: 0.0000
Epoch 15/15
94/94 [=====] - 1s 5ms/step - loss: 0.2001 - accuracy: 0.0000

```

accuracy=99.7 val_acc=97.5

```
import matplotlib.pyplot as plt
acc = history.history["accuracy"]
val_acc = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, "bo", label="Training accuracy")
plt.plot(epochs, val_acc, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
```





```
conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False)
conv_base.trainable = False

conv_base.trainable = True
print("This is the number of trainable weights "
      "before freezing the conv base:", len(conv_base.trainable_weights))

conv_base.trainable = False
print("This is the number of trainable weights "
      "after freezing the conv base:", len(conv_base.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 26
 This is the number of trainable weights after freezing the conv base: 0

✓ Utilizing data augmentation for feature extraction

```
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
```

```
import tensorflow as tf
from tensorflow import keras
```

```
!pip install tensorflow
```

```
Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: flatbuffers>=23.5.26 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: gast!=0.5.0,!=0.5.1,!=0.5.2,>=0.2.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: ml-dtypes~=0.2.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: numpy<2.0.0,>=1.23.5 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: protobuf!=4.21.0,!=4.21.1,!=4.21.2,!=4.21.3,!=4.21.4 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: wrapt<1.15,>=1.11.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: tensorboard<2.16,>=2.15 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: tensorflow-estimator<2.16,>=2.15.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: keras<2.16,>=2.15.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: google-auth-oauthlib<2,>=0.5 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: pyasn1<0.7.0,>=0.4.6 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.10/dist-packages
```

```
import tensorflow as tf
from tensorflow import keras
```

```
callbacks = [
```



```

keras.callbacks.ModelCheckpoint(
    filepath="feature_extraction_with_data_augmentationPT2.keras",
    save_best_only=True,
    monitor="val_loss",
    save_weights_only=True)
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

```

Epoch 1/30
94/94 [=====] - 16s 171ms/step - loss: 1.3165 - accu
Epoch 2/30
94/94 [=====] - 17s 174ms/step - loss: 1.7702 - accu
Epoch 3/30
94/94 [=====] - 16s 171ms/step - loss: 1.4329 - accu
Epoch 4/30
94/94 [=====] - 16s 166ms/step - loss: 1.3123 - accu
Epoch 5/30
94/94 [=====] - 16s 167ms/step - loss: 1.0921 - accu
Epoch 6/30
94/94 [=====] - 20s 213ms/step - loss: 0.9200 - accu
Epoch 7/30
94/94 [=====] - 16s 166ms/step - loss: 0.9962 - accu
Epoch 8/30
94/94 [=====] - 20s 210ms/step - loss: 0.8914 - accu
Epoch 9/30
94/94 [=====] - 16s 167ms/step - loss: 0.6181 - accu
Epoch 10/30
94/94 [=====] - 16s 171ms/step - loss: 0.7317 - accu
Epoch 11/30
94/94 [=====] - 16s 168ms/step - loss: 0.6439 - accu
Epoch 12/30
94/94 [=====] - 16s 169ms/step - loss: 0.8221 - accu
Epoch 13/30
94/94 [=====] - 16s 165ms/step - loss: 0.7171 - accu
Epoch 14/30
94/94 [=====] - 16s 167ms/step - loss: 0.5458 - accu
Epoch 15/30
94/94 [=====] - 16s 169ms/step - loss: 0.6724 - accu
Epoch 16/30
94/94 [=====] - 16s 167ms/step - loss: 0.5941 - accu
Epoch 17/30
94/94 [=====] - 16s 167ms/step - loss: 0.5272 - accu
Epoch 18/30
94/94 [=====] - 16s 166ms/step - loss: 0.4459 - accu
Epoch 19/30

```

```

94/94 [=====] - 16s 165ms/step - loss: 0.4504 - accu
Epoch 20/30
94/94 [=====] - 16s 168ms/step - loss: 0.6865 - accu
Epoch 21/30
94/94 [=====] - 20s 210ms/step - loss: 0.5540 - accu
Epoch 22/30
94/94 [=====] - 20s 207ms/step - loss: 0.6136 - accu
Epoch 23/30
94/94 [=====] - 16s 166ms/step - loss: 0.5023 - accu
Epoch 24/30
94/94 [=====] - 20s 210ms/step - loss: 0.4436 - accu
Epoch 25/30
94/94 [=====] - 16s 167ms/step - loss: 0.5005 - accu
Epoch 26/30
94/94 [=====] - 16s 166ms/step - loss: 0.4109 - accu
Epoch 27/30
94/94 [=====] - 20s 209ms/step - loss: 0.3969 - accu
Epoch 28/30
94/94 [=====] - 16s 166ms/step - loss: 0.5013 - accu
Epoch 29/30
94/94 [=====] - 20s 208ms/step - loss: 0.5413 - accu
Epoch 30/30
94/94 [=====] - 16s 167ms/step - loss: 0.6212 - accu

```

```

test_model = keras.models.load_model(
    "feature_extraction_with_data_augmentationPT2.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

```

32/32 [=====] - 3s 89ms/step - loss: 4.2593 - accuracy: 0.9690

Test accuracy: 0.969 accuracy=96.9% val_Acc=97.4% test_acc=97.4%

